

A decorative graphic consisting of a dark blue rectangle at the top, with thin dark blue lines extending from its corners. A diagonal line extends from the top-right corner towards the top-right of the slide. Another diagonal line extends from the bottom-left corner towards the bottom-left of the slide. The main title is centered below the rectangle.

Prediction of Boston House Price

**Presenters: Jinyi Shang
Yu Cao
Chaohui Li**

Contents

Part 01

Introduction

Part 02

Dataset Description

Part 03

Data Preprocessing

Part 04

Linear Regression

Part 05

Neural Network

Part 06

Ensemble method

Part 07

Summary

Part 08


Q&A

The slide features a minimalist design with dark blue geometric lines. A horizontal line at the top and a vertical line on the left form an L-shape. A horizontal line at the bottom and a vertical line on the right form another L-shape. Two diagonal lines, one in the top right and one in the bottom left, point towards the center. The text is centered within this frame.

PART 01

Introduction

Introduction



The house value is related to many factors. So our project will focus on solving the problem of predicting house price for house buyers and sellers.

The data is collected from Boston. There are total 506 entries and 13 features.

A decorative graphic consisting of dark blue lines. A horizontal line at the top and a vertical line on the left form an L-shape. A diagonal line extends from the top right towards the center. Another L-shape is formed by a horizontal line at the bottom and a vertical line on the right. A diagonal line extends from the bottom left towards the center.

PART 02

Dataset Description

Dataset Description

#	Column	Description	Non-Null Count	Dtype
0	CRIM	Crime rate per capita in towns	486 non-null	float64
1	ZN	Proportion of residential land	486 non-null	float64
2	INDUS	Proportion of non-commercial land in urban areas	486 non-null	float64
3	CHAS	Charles River Dummy Variable	486 non-null	float64
4	NOX	Environmental protection index	506 non-null	float64
5	RM	Number of rooms per house	506 non-null	float64
6	AGE	Proportion of self-occupied units built before 1940	486 non-null	float64
7	DIS	Weighted distance from Boston's five employment centers	506 non-null	float64
8	RAD	Convenience Index to Highway	506 non-null	int64
9	TAX	Real estate tax rate per US\$10,100	506 non-null	int64
10	PTRATIO	Teacher-student ratio in towns	506 non-null	float64
11	B	Proportion of blacks in towns	506 non-null	float64
12	LSTAT	Proportion of landlords belonging to the lower income class	486 non-null	float64
13	MEDV	Median of house price of self-occupied house	506 non-null	float64

Dataset Description

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
count	486	486	486	486	506	506	486
mean	3.612	11.212	11.084	0.070	0.555	6.285	68.519
std	8.720	23.389	6.836	0.255	0.116	0.703	28.000
min	0.006	0.000	0.460	0.000	0.385	3.561	2.900
median	0.254	0.000	9.690	0.000	0.538	6.209	76.800
max	88.976	100.000	27.740	1.000	0.871	8.780	100.000

	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
count	506	506	506	506	506	486	506
mean	3.795	9.549	408.237	18.456	356.674	12.715	22.533
std	2.106	8.707	168.537	2.165	91.295	7.1561	9.197
min	1.130	1.000	187.000	12.600	0.320	1.730	5.000
median	3.207	5.000	330.000	19.050	391.440	11.430	21.200
max	12.127	24.000	711.000	22.000	396.900	37.970	50.000

A decorative graphic consisting of several thin, dark blue lines. A horizontal line at the top left extends to the right, then turns 90 degrees down to form a vertical line. Another horizontal line at the bottom left extends to the right, then turns 90 degrees up to form a vertical line. These two vertical lines are connected by a horizontal line on the right side. Additionally, there are two diagonal lines: one in the top right corner and one in the bottom left corner, both extending from the outer edges towards the center.

PART 03

Data Preprocessing









The numbers of missing values for every feature:

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
TRUE	TRUE	TRUE	TRUE	FALSE	FALSE	TRUE
20	20	20	20	0	0	20

DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE
0	0	0	0	0	20	0



The distribution for features that exist missing values:





Three kinds of filling strategies:






```
imp1 = SimpleImputer(missing_values=np.nan, strategy='median')
imp2 = SimpleImputer(missing_values=np.nan, strategy='mean')
imp3 = SimpleImputer(missing_values=np.nan, strategy='most_frequent')

df = boston_housing.copy()
imp1.fit(np.array(df['CRIM']).reshape(-1, 1))
df['CRIM'] = imp1.transform(np.array(boston_housing['CRIM']).reshape(-1, 1))
imp1.fit(np.array(df['ZN']).reshape(-1, 1))
df['ZN'] = imp1.transform(np.array(boston_housing['ZN']).reshape(-1, 1))
imp2.fit(np.array(df['INDUS']).reshape(-1, 1))
df['INDUS'] = imp2.transform(np.array(boston_housing['INDUS']).reshape(-1, 1))
imp3.fit(np.array(df['CHAS']).reshape(-1, 1))
df['CHAS'] = imp3.transform(np.array(boston_housing['CHAS']).reshape(-1, 1))
imp1.fit(np.array(df['AGE']).reshape(-1, 1))
df['AGE'] = imp2.transform(np.array(boston_housing['AGE']).reshape(-1, 1))
imp1.fit(np.array(df['LSTAT']).reshape(-1, 1))
df['LSTAT'] = imp1.transform(np.array(boston_housing['LSTAT']).reshape(-1, 1))
```

Data Preprocessing



Transformation and Standardization



Non-linear
transformation



```
pt = preprocessing.PowerTransformer()  
features_pt = pt.fit_transform(features)
```

method='box-cox'
method='yeo-johnson'
Standardize=True

Box-Cox requires input data to be strictly positive.
Yeo-Johnson supports both positive and negative data.



`sklearn.feature_selection.SelectKBest(score_func, k)`

`score_func:`

For regression: `f_regression`, `mutual_info_regression`

`k:`


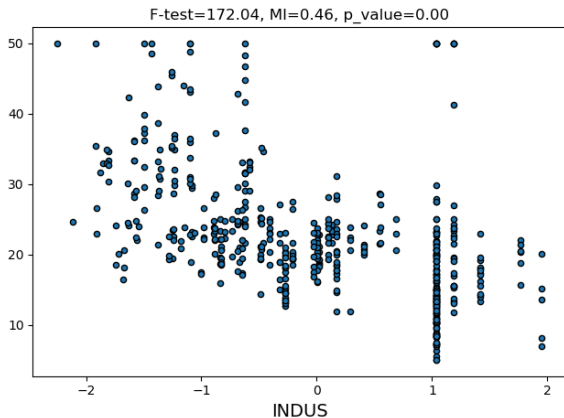


`k=int` or “all”

```
f_test, _ = f_regression(features_pt, price)
mi = mutual_info_regression(features_pt, price)
new = SelectKBest(f_regression, k='all')
new.fit_transform(features_pt, price)
print('p_values of features are:', new.pvalues_)
plt.figure(figsize=(15, 5))
for i in range(len(cols)):
    plt.scatter(features_pt[:, i], price, edgecolors='black', s=20)
    plt.xlabel("{}".format(cols[i]), fontsize=14)
    plt.title("F-test={:.2f}, MI={:.2f}, p_value={:.2f}".format(f_test[i], mi[i], new.pvalues_[i]))
plt.show()
```

Data Preprocessing









Feature Significance Test



Data Preprocessing






Feature Significance Test



Data Preprocessing



Feature Significance Test



summary of p_value

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
8.59E-24	2.08E-18	5.06E-34	3.17E-05	3.79E-24	4.33E-69	4.92E-18
DIS	RAD	TAX	PTRATIO	B	LSTAT	
9.35E-12	2.35E-14	1.32E-29	4.91E-35	7.22E-14	1.75E-107	

A decorative graphic consisting of a dark blue rectangular frame with a white diagonal line crossing it from the bottom-left to the top-right. The frame is composed of several segments: a horizontal line at the top, a vertical line on the left, a horizontal line at the bottom, and a vertical line on the right. A diagonal line segment crosses the frame from the bottom-left towards the top-right.

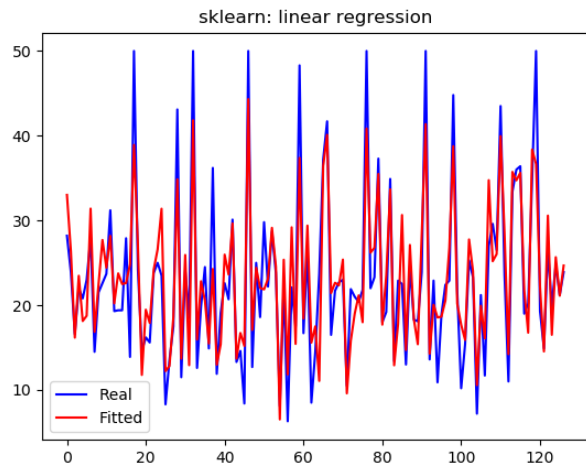
PART 04

Linear Regression

Results

MSE & R²: MSE: 20.239 R²: 0.796

Real vs. Fitted



```
# linear regression
lr1 = LinearRegression()
lr1.fit(x_train, y_train)
y_pred1 = lr1.predict(x_test)
plt.plot(range(len(y_test)), y_test, 'b', label='Real')
plt.plot(range(len(y_pred1)), y_pred1, 'r', label='Fitted')
plt.legend()
plt.title('sklearn: linear regression')
plt.savefig('Real vs. Fitted_LR.png')
plt.show()

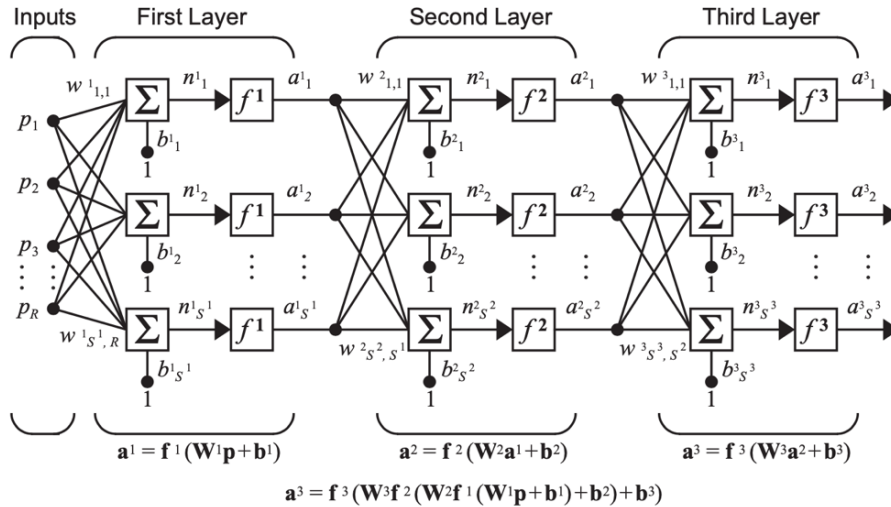
print('The coefficients of linear regression model are:', lr1.coef_)
print('The intercept of linear regression model is:', lr1.intercept_)
print("MSE:", metrics.mean_squared_error(y_test, y_pred1))
print("R^2:", r2_score(y_test, y_pred1))
```

A decorative graphic consisting of dark blue lines. A horizontal line at the top and a vertical line on the left form an L-shape. A second horizontal line is below the first, and a second vertical line is to the right of the first, creating a rectangular frame. Two diagonal lines are also present: one in the top right corner and one in the bottom left corner.

PART 05

Neural Network

Neural Network



Three-Layer Network

Neural Networks are a set of algorithms, modeled loosely on the human brain.

$$n = \mathbf{W}\mathbf{p} + \mathbf{b}$$

Net Input

$$a = f(\mathbf{W}\mathbf{p} + \mathbf{b})$$

Neuron Output

Parameters

Model:

MLPRegressor()

Parameters:

activation	activation function for the hidden layer
solver	the solver for weight optimization
batch_size	size of minibatches for stochastic optimizers
max_iter	the maximum number of iterations
early_stopping	whether to stop when validation score is not improving

```
def selepara_RSCV(Model, params, X, Y):
    n_iter = np.arange(200, 400, 50)

    score = 0

    for i in n_iter:
        clf = RandomizedSearchCV(Model, params, n_iter=i, n_jobs=-1, cv=5)

        grid_result = clf.fit(X, Y)

        if grid_result.best_score_ > score:
            score = grid_result.best_score_
            parameters = grid_result.best_params_
            n = i

    return n, score, parameters
```

```
# Parameters
hls = (100,)

optimizer = ['SGD', 'RMSprop', 'Adagrad', 'Adadelta', 'Adam', 'Adamax', 'Nadam']

param_random = dict(
    activation=['identity', 'logistic', 'tanh', 'relu'],
    solver=['lbfgs', 'sgd', 'adam'],
    max_iter=np.arange(7000, 15000, 1000),
    batch_size=np.arange(10, 200, 10),
    early_stopping=[False, True])

best_n_iter, best_score, paras = selepara_RSCV(NN, param_random, features, target)
```

♥ RandomizedSearchCV:

grid search for parameters by randomly sampling in the parameter space

Its search ability depends on the set 'n_iter' parameter.


♥ Best Parameters:

Parameters	activation	solver	batch_size	max_iter	early_stopping
Value	relu	sgd	70	11000	True


Results

MSE & R²: MSE: 10.352. R²: 0.896

Error vs. Epochs



Real vs. Fitted



```
# Train the model
nn = MLPRegressor(hidden_layer_sizes=hls,
                  activation=paras['activation'],
                  solver=paras['solver'],
                  max_iter=paras['max_iter'],
                  early_stopping=paras['early_stopping'],
                  batch_size=paras['batch_size'],
                  random_state=50)

MSE, R2, Loss_curve, prediction = train_model(nn)

print('MSE:', MSE)
print('R2:', R2)

print('Number of iteration when the model converges:', nn.n_iter_)

# Plot results

# Error and Epochs
pd.DataFrame(Loss_curve).plot()
plt.title('Error vs. Epochs')
plt.xlabel('Epochs')
plt.ylabel('MSE')
plt.legend('MSE')
plt.savefig('Error vs. Epochs.png')
plt.show()

# Real vs Fitted
plt.plot(range(len(y_test)), y_test, c='b')
plt.plot(range(len(prediction)), prediction, c='r')
plt.title('Real vs. Fitted')
plt.ylabel('Media Value of Home')
plt.legend(['Real', 'Fitted'])
plt.savefig('Real vs. Fitted.png')
plt.show()
```


A decorative graphic consisting of dark blue lines. A horizontal line at the top and a vertical line on the left form an L-shape. A horizontal line at the bottom and a vertical line on the right form another L-shape. Two diagonal lines, one in the top right and one in the bottom left, point towards the center.

PART 06

Ensemble method


Ensemble method

Reason:

Data size



Bad generalization



Ensemble method

01

Bagging

random forest

02

Boosing

adaboost, Gradient boosting




Ensemble method

01 Random forest

How to improve the ability of generalization?

Resampling technique(Bootstrap)

Randomly select subfeature



Ensemble method

Bootstrap:

data set D : m samples,

randomly take a sample from D
with replacement into set D'

data set D' : m samples.




Ensemble method

Randomly select subfeature

Traditional: k of features

RF : subfeature set

information gain



Ensemble method

RF Performance:

```
{'max_depth': 10, 'n_estimators': 100}
```


MSE: 9.3766947500555

R square: 0.905341820710926



Ensemble method

02 Gradient boosting



Given dataset:

$$T = \{(x_1, y_1), (x_2, y_2), \dots (x_N, y_N)\}, \quad y \in \{-1, +1\} \quad \text{and base estimator} \quad G_m(x)$$

initialize weight $w_i^{(1)} = \frac{1}{N}, \quad i = 1, 2, 3, \dots N$

for m=1 to M:

train G(x):
$$G_m(x) = \arg \min_{G(x)} \sum_{i=1}^N w_i^{(m)} \mathbb{I}(y_i \neq G(x_i))$$

Adaboost

calculate error rate

$$\epsilon_m = \frac{\sum_{i=1}^N w_i^{(m)} \mathbb{I}(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i^{(m)}}$$

update parameter


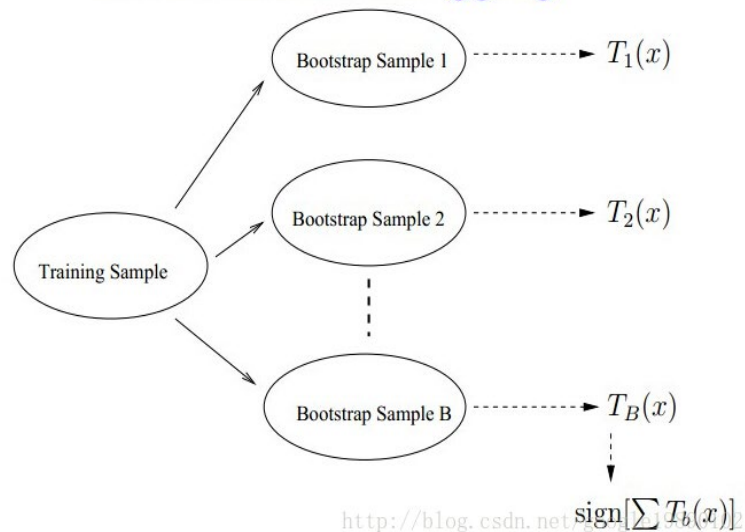
$$\alpha_m = \frac{1}{2} \ln \frac{1 - \epsilon_m}{\epsilon_m}$$
$$w_i^{(m+1)} = \frac{w_i^{(m)} e^{-y_i \alpha_m G_m(x_i)}}{Z^{(m)}}, \quad i = 1, 2, 3 \dots N$$

final model:


$$G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$$

Schematics

Schematics of Bagging



Schematics of Boosting



Ensemble method

GB Performance:

```
{'min_samples_leaf': 3, 'n_estimators': 150}
```

MSE: 9.126177702162598

R square: 0.9078708022194981



The slide features a minimalist design with dark blue geometric lines. A horizontal line at the top and a vertical line on the left form an L-shape. A horizontal line at the bottom and a vertical line on the right form another L-shape. Two diagonal lines, one in the top right and one in the bottom left, point towards the center. A dark blue rectangular bar is positioned in the upper center.

PART 07

Summary

Model	MSE	R square
Linear regression	20.239	0.796
Neural network	10.352	0.896
Random forest	9.376	0.905
Gradient boosting	9.126	0.907

Ensemble method is better.

A decorative graphic consisting of dark blue lines. A horizontal line at the top and a vertical line on the left form an L-shape. A horizontal line at the bottom and a vertical line on the right form another L-shape. Two diagonal lines, one in the top right and one in the bottom left, point towards the center.

PART 08

Q&A